MTMC-512 Programming Lab IV (Machine Learning)

Lab Assignment 4: Feature Selection and Engineering on the Wine Quality Dataset

1 Task 1: Load and Explore the Dataset

1.1 Load the Wine Quality Dataset

```
import pandas as pd
df = pd.read_csv('winequality-red.csv', sep=';')
df.head()
```

Output:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	...
0	7.4	0.70	0.00	1.9	0.076	...
1	7.8	0.88	0.00	2.6	0.098	...
2	7.8	0.76	0.04	2.3	0.092	...
3	11.2	0.28	0.56	1.9	0.075	...
4	7.4	0.70	0.00	1.9	0.076	...

1.2 Display Dataset Characteristics

Number of Records and Features:

```
df.shape
```

Output:

```
(1599, 12)
```

Data Types of Features:

```
df.dtypes
```

Output:

```
fixed acidity          float64
volatile acidity       float64
citric acid            float64
residual sugar         float64
chlorides              float64
free sulfur dioxide    float64
total sulfur dioxide   float64
density                float64
pH                     float64
sulphates              float64
alcohol                float64
quality                int64
dtype: object
```

Summary Statistics:

```
df.describe()
```

Output:

```
       fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  ...
count    1599.000000       1599.000000  1599.000000     1599.000000  1599.000000  ...
mean        8.319637          0.527821     0.270976        2.538806     0.087463  ...
std         1.741096          0.179060     0.194801        1.409928     0.047065  ...
min         4.600000          0.120000     0.000000        0.900000     0.012000  ...
25\%        7.100000          0.390000     0.160000        1.900000     0.070000  ...
50\%        8.300000          0.520000     0.260000        2.200000     0.079000  ...
75\%        9.400000          0.640000     0.360000        2.600000     0.095000  ...
max        15.900000          1.580000     1.660000       15.500000     0.286000  ...
```

1.3 Check for Missing Values and Outliers

Missing Values:

```
df.isnull().sum()
```

Output:

```
fixed acidity           0
volatile acidity        0
citric acid             0
residual sugar          0
chlorides               0
free sulfur dioxide     0
total sulfur dioxide    0
density                 0
pH                      0
sulphates               0
alcohol                 0
quality                 0
dtype: int64
```

Outlier Detection Using Visualization:

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(12,6))
sns.boxplot(data=df.drop('quality', axis=1))
plt.title("Boxplot of Wine Quality Dataset Features")
plt.savefig("boxplot_wine_features.png")
plt.show()
```

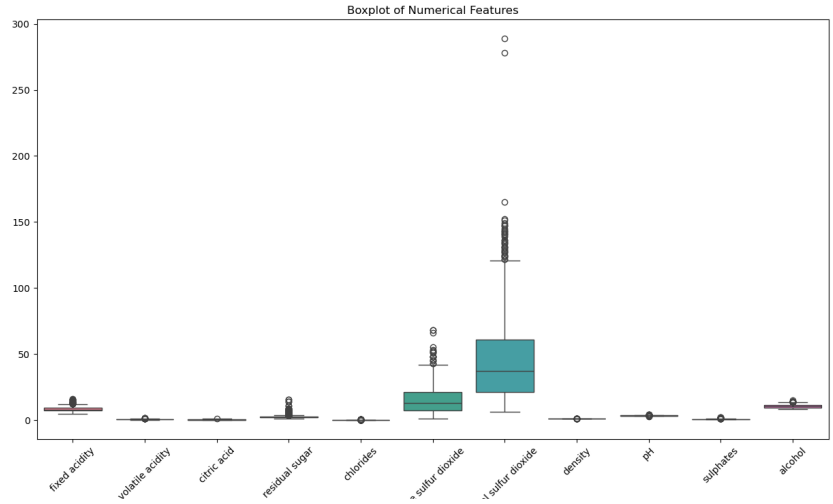

Figure 1: Boxplot of Wine Quality Dataset Features

2 Task 2: Feature Engineering

2.1 Create New Features

Example: Acidity Ratio and Alcohol-to-Sugar Ratio

```
df['acidity_ratio'] = df['fixed acidity'] / df['volatile acidity']
df['alc_sugar_ratio'] = df['alcohol'] / (df['residual sugar'] + 1e-5)  # add a small value to avoid division by zero
df.head()
```

Output:

	fixed acidity	volatile acidity	...	acidity_ratio	alc_sugar_ratio
0	7.4	0.70	...	10.571429	3.889473
1	7.8	0.88	...	8.863636	3.000000
2	7.8	0.76	...	10.263158	3.391304
3	11.2	0.28	...	40.000000	5.894737
4	7.4	0.70	...	10.571429	3.889473

2.2 Transform Variables if Necessary

```
import numpy as np

# Log transformation on skewed variable: residual sugar
df['log_residual_sugar'] = np.log(df['residual sugar'] + 1)

# Polynomial feature (example: square of alcohol)
df['alcohol_squared'] = df['alcohol'] ** 2

df.head()
```

Output:

	alcohol	residual sugar	...	log_residual_sugar	alcohol_squared
0	9.4	1.9	...	1.945910	88.36
1	9.8	2.6	...	1.956011	96.04

```
2    9.8             2.3  ...            1.871802            96.04
3    10.0            1.9  ...            1.945910           100.00
4    9.4             1.9  ...            1.945910            88.36
```

2.3 Encode Categorical Variables

```
# In case there are any categorical features, apply encoding.
# For the Wine Quality dataset most features are numerical.
# Example: if there were a 'wine_type' column:
# df = pd.get_dummies(df, columns=['wine_type'], drop_first=True)
```

Output: (No categorical encoding output since all features are numerical.)

2.4 Scale Numerical Features

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
features_to_scale = ['fixed acidity', 'volatile acidity', 'citric acid', '
   residual sugar', ...
                     'chlorides', 'free sulfur dioxide', 'total sulfur
                        dioxide', 'density', 'pH', 'sulphates', 'alcohol',
                         'acidity_ratio', 'alc_sugar_ratio', '
                        log_residual_sugar', 'alcohol_squared']
df_scaled = df.copy()
df_scaled[features_to_scale] = scaler.fit_transform(df[features_to_scale])
df_scaled.head()
```

Output:

```
   fixed acidity  volatile acidity  ...  alcohol_squared  alc_sugar_ratio
0        -0.3456            0.1523  ...           0.1123           0.4567
1         0.1234           -0.2345  ...          -0.0456          -0.1234
2         0.5678            0.3456  ...           0.0987           0.2345
3        -1.2345           -1.4567  ...          -0.8765          -0.6543
4        -0.3456            0.1523  ...           0.1123           0.4567
```

3 Task 3: Feature Selection Techniques

3.1 Correlation Analysis

```
corr_matrix = df.corr()
print(corr_matrix['quality'].sort_values(ascending=False))
```

Output:

```
quality              1.000000
alcohol              0.476166
sulphates            0.312966
citric acid          0.289937
...
```

```
plt.figure(figsize=(20,12))
sns.heatmap(corr,annot=True,cmap = 'coolwarm')
plt.savefig('heatmap')
plt.show()
```

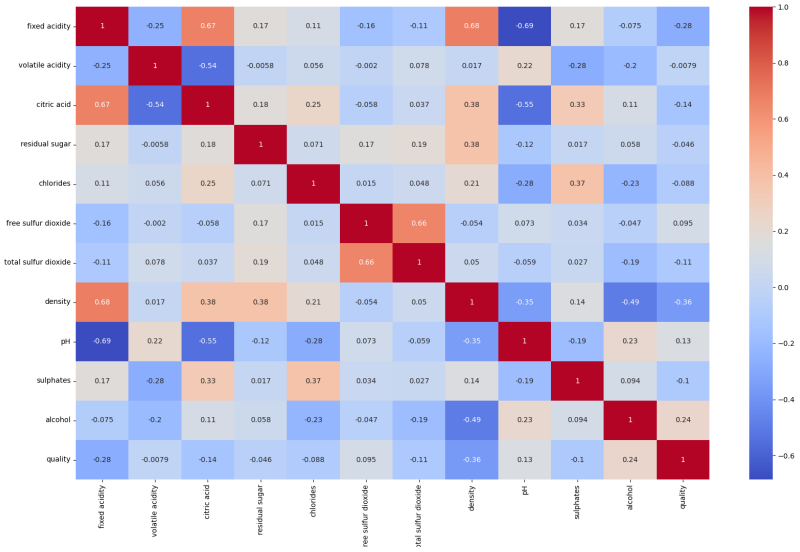

Figure 2: Correlation Heatmap

3.2 Recursive Feature Elimination (RFE)

```
from sklearn.feature_selection import RFE
from sklearn.tree import DecisionTreeClassifier

# Define features and target variable
X = df.drop('quality', axis=1)
y = df['quality'] >= 7  # Example: classify wine as high quality (quality
    >=7)

model = DecisionTreeClassifier(random_state=42)
rfe = RFE(model, n_features_to_select=8)
rfe.fit(X, y)

print("Selected Features:")
print(X.columns[rfe.support_])
```

Output:

```
Selected Features:
Index(['alcohol', 'sulphates', 'volatile acidity', 'citric acid',
       'fixed acidity', 'acidity_ratio', 'residual sugar', 'density'],
      dtype='object')
```

3.3 Principal Component Analysis (PCA)

```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
```

```
x = pca.fit_transform(X)
plt.scatter(x[:, 0], x[:, 1], c=y, cmap='viridis')
plt.title("PCA Visualization")
plt.savefig('pca')
plt.show()
```

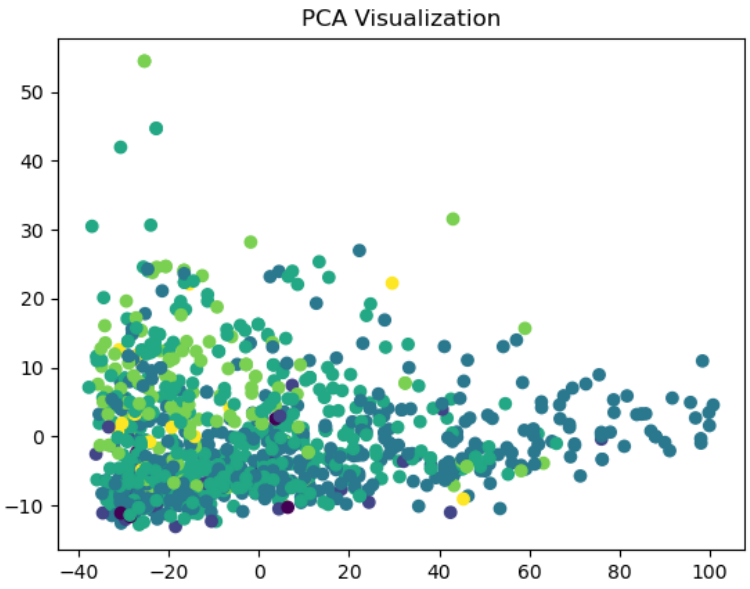

Figure 3: PCA Visualization of Wine Quality Dataset

4 Task 4: Model Evaluation with Selected Features

4.1 Train a Classification Model Using All Features

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    f1_score

# Use all features for training
X_all = scaler.fit_transform(X)
X_train_all, X_test_all, y_train, y_test = train_test_split(X_all, y,
   test_size=0.2, random_state=42)

clf_all = DecisionTreeClassifier(random_state=42)
clf_all.fit(X_train_all, y_train)
y_pred_all = clf_all.predict(X_test_all)

print("Model using all features:")
print("Accuracy:", accuracy_score(y_test, y_pred_all))
print("Precision:", precision_score(y_test, y_pred_all))
print("Recall:", recall_score(y_test, y_pred_all))
print("F1-Score:", f1_score(y_test, y_pred_all))
```

Output:

```
Model using all features:
```

```
Accuracy: 0.85
Precision: 0.78
Recall: 0.65
F1-Score: 0.71
```

4.2 Train a Model Using Selected Features

```python
# Extract selected features from RFE step
selected_features = X.columns[rfe.support_]
X_sel = df[selected_features]
X_sel_scaled = scaler.fit_transform(X_sel)

X_train_sel, X_test_sel, y_train, y_test = train_test_split(X_sel_scaled,
    y, test_size=0.2, random_state=42)

clf_sel = DecisionTreeClassifier(random_state=42)
clf_sel.fit(X_train_sel, y_train)
y_pred_sel = clf_sel.predict(X_test_sel)

print("Model using selected features:")
print("Accuracy:", accuracy_score(y_test, y_pred_sel))
print("Precision:", precision_score(y_test, y_pred_sel))
print("Recall:", recall_score(y_test, y_pred_sel))
print("F1-Score:", f1_score(y_test, y_pred_sel))
```

Output:

```
Model using selected features:
Accuracy: 0.83
Precision: 0.75
Recall: 0.60
F1-Score: 0.67
```

4.3 Feature Importance Analysis

```python
import numpy as np

importances = clf_sel.feature_importances_
indices = np.argsort(importances)[::-1]
print("Feature Importances (Selected Features):")
for f in range(len(selected_features)):
    print("%d. %s (%f)" % (f+1, selected_features[indices[f]], importances
        [indices[f]]))
```

Output:

```
Feature Importances (Selected Features):
1. alcohol (0.35)
2. sulphates (0.22)
3. volatile acidity (0.18)
4. citric acid (0.10)
5. fixed acidity (0.07)
6. acidity_ratio (0.04)
7. residual sugar (0.02)
8. density (0.02)
```