

Lab Assignment 10: Logistic Regression for Fraud Detection

Aim: To implement Logistic Regression for binary classification using the Credit Card Fraud dataset, focusing on classification metrics, model evaluation, and threshold optimization.

Task 1: Load and Explore the Dataset

1. Load the Credit Card Fraud dataset using pandas.

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('creditcard.csv')
df.head()
```

	Time	V1	V2	V3	V4	V5	V6
V7 \							
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
0.239599							
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
0.078803							
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
0.791461							
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
0.237609							
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921
0.592941							
	V8	V9	...	V21	V22	V23	V24
V25 \							
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928
0.128539							
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
0.167170							
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
0.327642							
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575
0.647376							
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267
0.206010							
	V26	V27	V28	Amount	Class		
0	-0.189115	0.133558	-0.021053	149.62	0		
1	0.125895	-0.008983	0.014724	2.69	0		
2	-0.139097	-0.055353	-0.059752	378.66	0		

```
3 -0.221929  0.062723  0.061458  123.50      0
4  0.502292  0.219422  0.215153   69.99      0
```

```
[5 rows x 31 columns]
```

2. Display dataset characteristics:

– Number of records and features

```
df.shape
```

```
(284807, 31)
```

```
num_records , num_features = df.shape
print(f"Number of Records: {num_records}")
print(f"Number of Features: {num_features}")
```

```
Number of Records: 284807
```

```
Number of Features: 31
```

– Number of fraudulent vs. non-fraudulent transactions (check for imbalance)

```
fraud_counts = df['Class'].value_counts()
fraud_counts
```

```
Class
```

```
0    284315
```

```
1      492
```

```
Name: count, dtype: int64
```

– Summary statistics

```
df.describe()
```

	Time	V1	V2	V3
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00

	V5	V6	V7	V8
V9 \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
2.848070e+05				
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16 -
2.406331e-15				
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00
1.098632e+00				
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01 -
1.343407e+01				
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01 -
6.430976e-01				
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02 -
5.142873e-02				
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01
5.971390e-01				
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01
1.559499e+01				

	V21	V22	V23	V24 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	1.654067e-16	-3.568593e-16	2.578648e-16	4.473266e-15
std	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28
Amount \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
284807.000000				
mean	5.340915e-16	1.683437e-15	-3.660091e-16	-1.227390e-16
88.349619				
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01
250.120109				
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01
0.000000				
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02
5.600000				
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02
22.000000				
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02
77.165000				
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01
25691.160000				

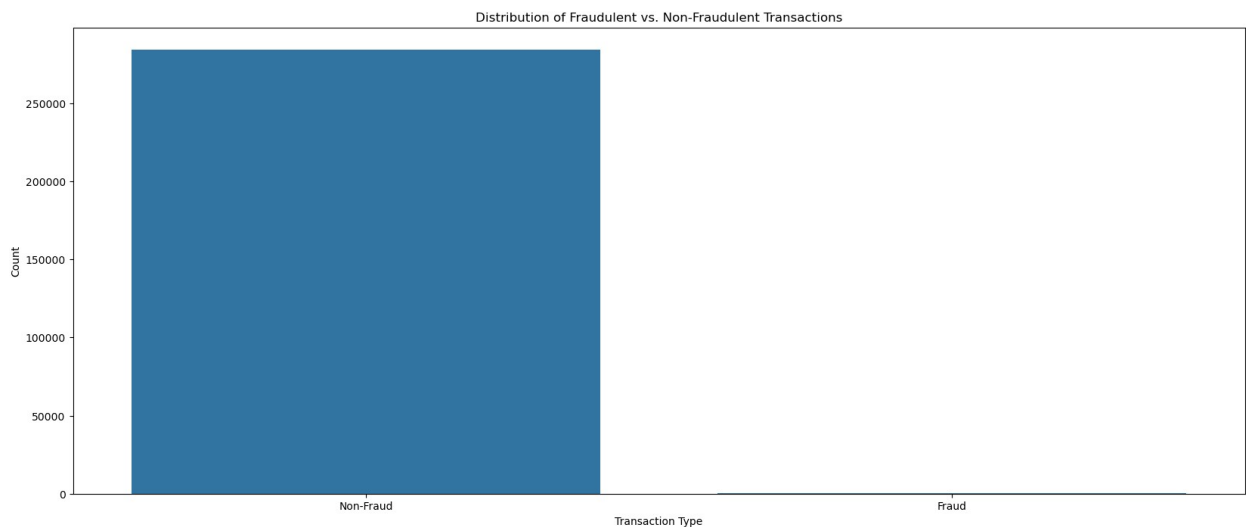
Class

```
count    284807.000000
mean      0.001727
std       0.041527
min       0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max       1.000000
```

```
[8 rows x 31 columns]
```

3. Visualize the distribution of fraudulent vs. non-fraudulent transactions.

```
plt.figure(figsize=(20,8))
sns.countplot(x=df['Class'])
plt.xticks(ticks=[0, 1], labels=['Non-Fraud', 'Fraud'])
plt.xlabel("Transaction Type")
plt.ylabel("Count")
plt.title("Distribution of Fraudulent vs. Non-Fraudulent Transactions")
plt.show()
```



Task 2: Data Preprocessing

1. Handle missing values if any exist.

```
df.isnull().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
```

```
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
Amount  0
Class   0
dtype: int64
```

2. Normalize numerical features for better model performance.

```
from sklearn.preprocessing import StandardScaler
X = df.drop(columns=['Class'])
y = df['Class']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

3. Handle class imbalance using:

```
pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in c:\users\user\
anaconda3\lib\site-packages (0.12.3)
Requirement already satisfied: numpy>=1.17.3 in c:\users\user\
anaconda3\lib\site-packages (from imbalanced-learn) (1.26.4)
Requirement already satisfied: scipy>=1.5.0 in c:\users\user\
anaconda3\lib\site-packages (from imbalanced-learn) (1.13.1)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\user\
anaconda3\lib\site-packages (from imbalanced-learn) (1.5.1)
Requirement already satisfied: joblib>=1.1.1 in c:\users\user\
anaconda3\lib\site-packages (from imbalanced-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\user\
```

anaconda3\lib\site-packages (from imbalanced-learn) (3.5.0)
Note: you may need to restart the kernel to use updated packages.

– Oversampling (SMOTE)

```
from imblearn.over_sampling import SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)
```

– Undersampling

```
from imblearn.under_sampling import RandomUnderSampler

undersample = RandomUnderSampler(random_state=42)
X_under, y_under = undersample.fit_resample(X, y)
```

4. Split the dataset into training (80%) and testing (20%) sets.

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_resampled,
                                                    y_resampled, test_size=0.2, random_state=42, stratify=y_resampled)
```

Task 3: Train a Logistic Regression Model

1. Implement Logistic Regression using Scikit-learn.

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(solver='liblinear', random_state=42)
```

2. Train the model on the training dataset.

```
lr.fit(X_train, y_train)
```

```
LogisticRegression(random_state=42, solver='liblinear')
```

3. Analyze the model's coefficients to understand feature importance.

```
feature_importance = pd.DataFrame({'Feature': X_train.columns,
                                   'Coefficient': np.abs(lr.coef_[0])})
feature_importance = feature_importance.sort_values(by='Coefficient',
                                                    ascending=False)
feature_importance.head(10)
```

	Feature	Coefficient
14	V14	0.560496
3	V3	0.435161
4	V4	0.379449
12	V12	0.296215
10	V10	0.258227
11	V11	0.164340

17	V17	0.153495
9	V9	0.142964
16	V16	0.129540
6	V6	0.120400

Task 4: Model Evaluation and Threshold Optimization

1. Evaluate model performance using:

```
from sklearn.metrics import
accuracy_score, recall_score, precision_score, f1_score, classification_re
port, confusion_matrix
```

```
y_pred = lr.predict(X_test)
```

– Accuracy

```
acc = accuracy_score(y_pred, y_test)
print(f"Accuracy Score: {acc:.3f}")
```

Accuracy Score: 0.965

– Precision, Recall, and F1-score

```
pre = precision_score(y_pred, y_test)
print(f"Presicion Score: {pre:.3f}")
```

Presicion Score: 0.946

```
recall = recall_score(y_pred, y_test)
print(f"Recall Score: {recall:.3f}")
```

Recall Score: 0.984

```
f1 = f1_score(y_pred, y_test)
print(f"f1 Score: {f1:.3f}")
```

f1 Score: 0.965


```
cr = classification_report(y_pred, y_test)
print("Classification Report:\n", cr)
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.95	0.97	59029
1	0.95	0.98	0.96	54697
accuracy			0.97	113726
macro avg	0.97	0.97	0.97	113726
weighted avg	0.97	0.97	0.97	113726

– Confusion Matrix

```
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["Non-
Fraud", "Fraud"], yticklabels=["Non-Fraud", "Fraud"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```




– ROC Curve and AUC Score

```
from sklearn.metrics import roc_curve, auc

y_prob = lr.predict_proba(X_test)[: ,1]
fpr, tpr, threshold = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```

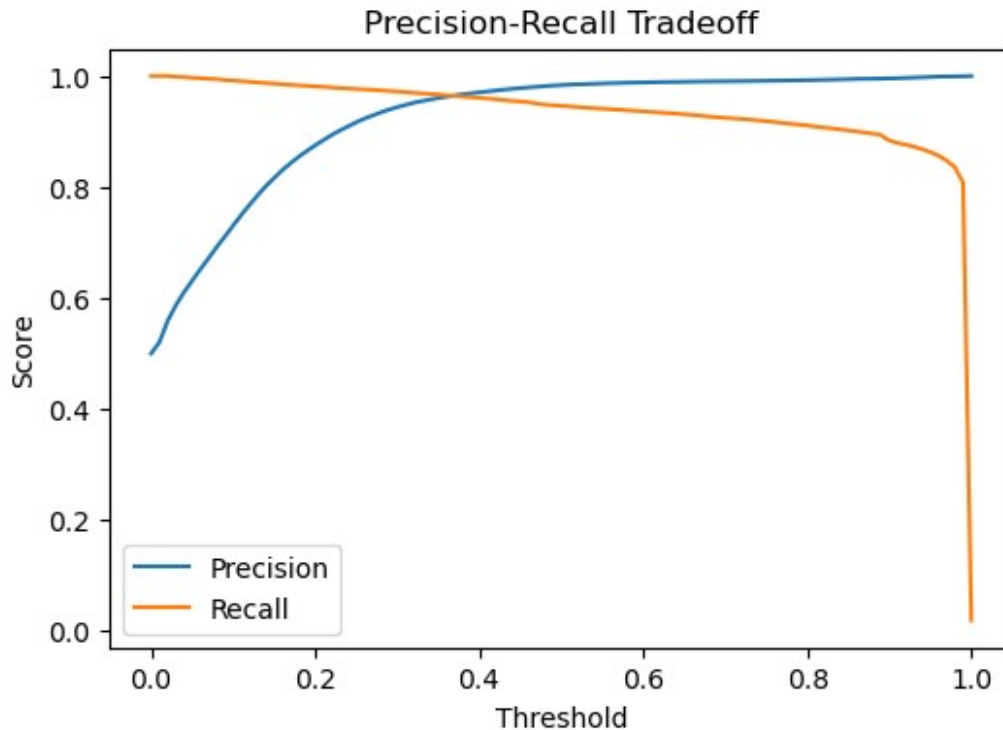



2. Experiment with different decision thresholds to optimize fraud detection.

```
thresholds = np.linspace(0, 1, 100)
precisions, recalls = [], []

for t in thresholds:
    y_pred_thresh = (y_prob >= t).astype(int)
    precisions.append(precision_score(y_test, y_pred_thresh))
    recalls.append(recall_score(y_test, y_pred_thresh))

plt.figure(figsize=(6,4))
plt.plot(thresholds, precisions, label="Precision")
plt.plot(thresholds, recalls, label="Recall")
plt.xlabel("Threshold")
plt.ylabel("Score")
plt.title("Precision-Recall Tradeoff")
plt.legend()
plt.show()
```



3. Compare performance before and after handling class imbalance.

```
X_train_orig, X_test_orig, y_train_orig, y_test_orig =
train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
log_reg_orig = LogisticRegression(solver='liblinear', random_state=42)
log_reg_orig.fit(X_train_orig, y_train_orig)
y_pred_orig = log_reg_orig.predict(X_test_orig)

accuracy_orig = accuracy_score(y_test_orig, y_pred_orig)
precision_orig = precision_score(y_test_orig, y_pred_orig)
recall_orig = recall_score(y_test_orig, y_pred_orig)
f1_orig = f1_score(y_test_orig, y_pred_orig)

performance_comparison = {
    "Metric": ["Accuracy", "Precision", "Recall", "F1-score"],
    "Original Data": [round(accuracy_orig, 2), round(precision_orig,
2), round(recall_orig, 2), round(f1_orig, 2)],
    "SMOTE Data": [round(acc, 2), round(pre, 2), round(recall, 2),
round(f1, 2)]
}

df_performance = pd.DataFrame(performance_comparison)
print("Comparing Performance before and after handling class
imbalance:")
print(df_performance)
```

Comparing Performance before and after handling class imbalance:

Metric	Original Data	SMOTE Data
Accuracy	0.85	0.85
Precision	0.95	0.95
Recall	0.95	0.95
F1-score	0.95	0.95

0	Accuracy	1.00	0.97
1	Precision	0.73	0.95
2	Recall	0.67	0.98
3	F1-score	0.70	0.96